


An individual benchmark instance is a complete evidence record

1. Stimulus shown to the model



2. Standardized elicitation

Estimate a , b , c , and d from the item characteristic curve.

Return JSON with the requested fields and a short rationale.

3. Structured observation and scoring

```
{"a": 1.42, "b": -0.36,  
 "c": 0.18, "d": 1.00}
```

Compare each field with stored ground truth.
Retain parsing failures, rationale score,
model version, prompt, and latency.